


FT09 runtime story

main.py to persisted trace, with the real feedback loops highlighted

The viewer is built from what this runtime actually stores: recording frames, GoalBoard summaries, and run-log reflexion lines.



What the visualizer can show faithfully

- 1. action and state change from recording frames
- 2. chosen hypothesis, verdict, and final skill memory from GoalBoard
- 3. reflexion directives from run logs when they were persisted

What is still missing in current persistence

- raw M1 candidate set for each turn
- full raw M2 JSON including expected_step_diffs
- complete semantic-packet snapshot per turn

Code anchors: main.py, agents/swarm.py, agents/templates/agentica_simple/agent.py, goal_board.py, reflexion_prompt.py